

Course 395: Machine Learning

- Lecturers: Maja Pantic (maja@doc.ic.ac.uk)
Stephen Muggleton (shm@doc.ic.ac.uk)
- Goal (Lectures): To present basic theoretical concepts and key algorithms that form the core of machine learning
- Goal (CBC): To enable hands-on experience with implementing machine learning algorithms using Matlab
- Material: *Machine Learning* by Tom Mitchell (1997)
Manual for completing the CBC
Syllabus on CBR
Notes on Inductive Logic Programming
- More Info: <https://www.ibug.doc.ic.ac.uk/courses>

Course 395: Machine Learning – Lectures

- Lecture 1-2: Concept Learning (*M. Pantic*)
- Lecture 3-4: Decision Trees & CBC Intro (*M. Pantic*)
- Lecture 5-6: Artificial Neural Networks (*THs*)
- Lecture 7-8: Instance Based Learning (*M. Pantic*)
- Lecture 9-10: Genetic Algorithms (*M. Pantic*)
- Lecture 11-12: Evaluating Hypotheses (*THs*)
- Lecture 13-14: Guest Lectures on ML Applications
- Lecture 15-16: Inductive Logic Programming (*S. Muggleton*)
- Lecture 17-18: Inductive Logic Programming (*S. Muggleton*)

Course 395: Machine Learning – Exam Material

- Lecture 1-2: Concept Learning (*Mitchell*: Ch.1, Ch.2)
- Lecture 3-4: Decision Trees & CBC Intro (*Mitchell*: Ch.3)
- Lecture 5-6: Artificial Neural Networks (*Mitchell*: Ch.4)
- Lecture 7-8: Instance Based Learning (Syllabus, *Mitchell*: Ch.8)
- Lecture 9-10: Genetic Algorithms (*Mitchell*: Ch.9)
- Lecture 11-12: Evaluating Hypotheses (*Mitchell*: Ch.5)
- Lecture 13-14: not examinable
- Lecture 15-16: Inductive Logic Programming (*Notes*)
- Lecture 17-18: Inductive Logic Programming (*Notes*)

Course 395: Machine Learning - CBC

- Lecture 1-2: Concept Learning
- ➡ Lecture 3-4: Decision Trees & CBC Intro
- ➡ Lecture 5-6: Artificial Neural Networks
- ➡ Lecture 7-8: Instance Based Learning
- Lecture 9-10: Genetic Algorithms
- ➡ Lecture 11-12: Evaluating Hypotheses
- Lecture 13-14: Guest Lectures on ML Applications
- Lecture 15-16: Inductive Logic Programming
- Lecture 17-18: Inductive Logic Programming

Course 395: Machine Learning

NOTE

CBC accounts for 33% of the final grade for the Machine Learning Exam.

$$\text{final grade} = 0.66 * \text{exam_grade} + 0.33 * \text{CBC_grade}$$

Course 395: Machine Learning - CBC

- Lecture 1-2: Concept Learning
- ➡ Lecture 3-4: Decision Trees & **CBC Intro**
- ➡ Lecture 5-6: Artificial Neural Networks
- ➡ Lecture 7-8: Instance Based Learning
- Lecture 9-10: Genetic Algorithms
- ➡ Lecture 11-12: Evaluating Hypotheses
- Lecture 13-14: Guest Lectures on ML Applications
- Lecture 15-16: Inductive Logic Programming
- Lecture 17-18: Inductive Logic Programming

Course 395: Machine Learning – Lectures

➤ Lecture 1-2: Concept Learning (*M. Pantic*)

- Lecture 3-4: Decision Trees & CBC Intro (*M. Pantic*)
- Lecture 5-6: Artificial Neural Networks (*THs*)
- Lecture 7-8: Instance Based Learning (*M. Pantic*)
- Lecture 9-10: Genetic Algorithms (*M. Pantic*)
- Lecture 11-12: Evaluating Hypotheses (*THs*)
- Lecture 13-14: Guest Lectures on ML Applications
- Lecture 15-16: Inductive Logic Programming (*S. Muggleton*)
- Lecture 17-18: Inductive Logic Programming (*S. Muggleton*)

Concept Learning – Lecture Overview

- Why machine learning?
- Well-posed learning problems
- Designing a machine learning system
- Concept learning task
- Concept learning as Search
- Find-S algorithm
- Candidate-Elimination algorithm

Machine Learning

- Learning \leftrightarrow Intelligence
(Def: *Intelligence is the ability to learn and use concepts to solve problems.*)
- Machine Learning \leftrightarrow Artificial Intelligence
 - Def: *AI is the science of making machines do things that require intelligence if done by men* (Minsky 1986)
 - Def: *Machine Learning is an area of AI concerned with development of techniques which allow machines to learn*
- Why Machine Learning? \leftrightarrow Why Artificial Intelligence?

Machine Learning



Machine Learning

- Learning \leftrightarrow Intelligence
(Def: *Intelligence is the ability to learn and use concepts to solve problems.*)
- Machine Learning \leftrightarrow Artificial Intelligence
 - Def: *AI is the science of making machines do things that require intelligence if done by men* (Minsky 1986)
 - Def: *Machine Learning is an area of AI concerned with development of techniques which allow machines to learn*
- Why Machine Learning? \leftrightarrow Why Artificial Intelligence?
 - \equiv To build machines exhibiting intelligent behaviour (i.e., able to reason, predict, and adapt) while helping humans work, study, and entertain themselves

Machine Learning

- Machine Learning \leftrightarrow Artificial Intelligence
- Machine Learning \leftarrow Biology (e.g., Neural Networks, Genetic Algorithms)
- Machine Learning \leftarrow Cognitive Sciences (e.g., Case-based Reasoning)
- Machine Learning \leftarrow Statistics (e.g., Support Vector Machines)
- Machine Learning \leftarrow Probability Theory (e.g., Bayesian Networks)
- Machine Learning \leftarrow Logic (e.g., Inductive Logic Programming)
- Machine Learning \leftarrow Information Theory (e.g., used by Decision Trees)

Machine Learning

- Human Learning \leftrightarrow Machine Learning
 - human-logic inspired problem solvers (e.g., rule-based reasoning)
 - biologically inspired problem solvers (e.g., Neural Networks)
 - supervised learning - generates a function that maps inputs to desired outputs
 - unsupervised learning - models a set of inputs, labelled examples are not available
 - learning by education (e.g., reinforcement learning, case-based reasoning)
- General Problem Solvers vs. Purposeful Problem Solvers
 - emulating general-purpose human-like problem solving is impractical
 - restricting the problem domain results in ‘rational’ problem solving
 - example of General Problem Solver: Turing Test
 - examples of Purposeful Problem Solvers: speech recognisers, face recognisers, facial expression recognisers, data mining, games, etc.
- Application domains: security, medicine, education, finances, genetics, etc.

Well-posed Learning Problems

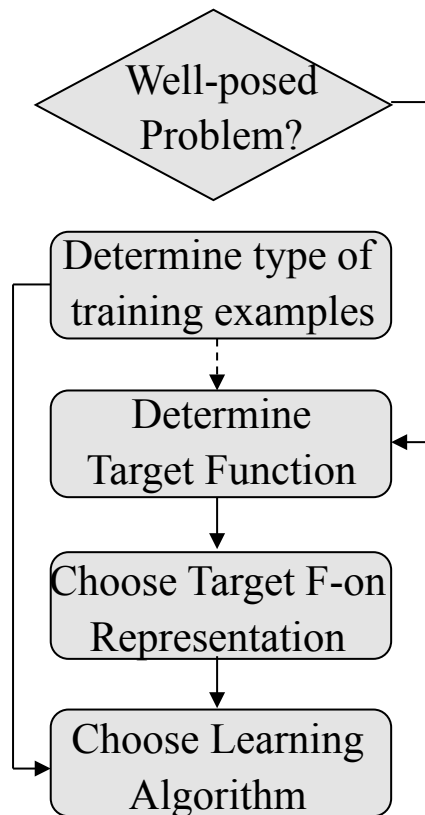
- Def 1 (Mitchell 1997):

A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves by experience E .

- Def 2 (Hadamard 1902):


A (machine learning) problem is well-posed if a solution to it exists, if that solution is unique, and if that solution depends on the data / experience but it is not sensitive to (reasonably small) changes in the data / experience.

Designing a Machine Learning System



- Target Function V represents the problem to be solved (e.g., choosing the best next move in chess, identifying people, classifying facial expressions into emotion categories)
- $V: D \rightarrow C$ where D is the input state space and C is the set of classes
 $V: D \rightarrow [-1, 1]$ is a general target function of a binary classifier
- Ideal Target Function is usually not known; machine learning algorithms learn an approximation of V , say V'
- Representation of function V' to be learned should
 - be as close an approximation of V as possible
 - require (reasonably) small amount of training data to be learned
- $V'(d) = w_0 + w_1x_1 + \dots + w_nx_n$ where $\langle x_1 \dots x_n \rangle \equiv d \in D$ is an input state. This reduces the problem to learning (the most optimal) weights w .

Designing a Machine Learning System



- $V: D \rightarrow C$ where D is the input state and C is the set of classes
 $V: D \rightarrow [-1, 1]$ is a general target function of a binary classifier
- $V'(d) = w_0 + w_1x_1 + \dots + w_nx_n$ where $\langle x_1 \dots x_n \rangle \equiv d \in D$ is an input state. This reduces the problem to learning (the most optimal) weights w .
- Training examples suitable for the given target function representation V' are pairs $\langle d, c \rangle$ where $c \in C$ is the desired output (classification) of the input state $d \in D$.
- Learning algorithm learns the most optimal set of weights w (so-called *best hypothesis*), i.e., the set of weights that best fit the training examples $\langle d, c \rangle$.
- Learning algorithm is selected based on the availability of training examples (supervised vs. unsupervised), knowledge of the final set of classes C (offline vs. online, i.e., eager vs. lazy), availability of a tutor (reinforcement learning).
- The learned V' is then used to solve new instances of the problem.

Concept Learning

- Concept learning
 - supervised, eager learning
 - target problem: whether something belongs to the target concept or not
 - target function: $V: D \rightarrow \{\text{true}, \text{false}\}$
- Underlying idea: Humans acquire general concepts from specific examples (e.g., concepts: living organism, beauty, computer, well-fitting-shoes) (note: each concept can be thought of as Boolean-valued function)
- Concept learning is inferring a Boolean-valued function from training data
→ concept learning is the prototype binary classification

Concept Learning Task – An Example

- Concept learning task:
 - target concept: Girls who Simon likes
 - target function: $c: D \rightarrow \{0, 1\}$
 - data $d \in D$: Girls, each described in terms of the following attributes
 - $a_1 \equiv \text{Hair}$ (possible values: blond, brown, black)
 - $a_2 \equiv \text{Body}$ (possible values: thin, normal, plump)
 - $a_3 \equiv \text{likesSimon}$ (possible values: yes, no)
 - $a_4 \equiv \text{Pose}$ (possible values: arrogant, natural, goofy)
 - $a_5 \equiv \text{Smile}$ (possible values: none, pleasant, toothy)
 - $a_6 \equiv \text{Smart}$ (possible values: yes, no)
 - target f-on representation: $h \equiv c': \langle a_1, a_2, a_3, a_4, a_5, a_6 \rangle \rightarrow \{0, 1\}$
 - training examples D : positive and negative examples of target function c
- **Aim:** Find a hypothesis $h \in H$ such that $(\forall d \in D) h(d) - c(d) < \varepsilon \approx 0$, where H is the set of all possible hypotheses $h \equiv \langle a_1, a_2, a_3, a_4, a_5, a_6 \rangle$, where each a_k , $k = [1..6]$, may be '?' (\equiv any value is acceptable), '0' (\equiv no value is acceptable), or a specific value.

Concept Learning Task – Notation

- Concept learning task:
 - target concept: Girls who Simon likes
 - target function: $c: D \rightarrow \{0, 1\}$
 - data $d \in D$: Girls, each described in terms of the following attributes
 - $a_1 \equiv \text{Hair}$ (possible values: blond, brown, black)
 - $a_2 \equiv \text{Body}$ (possible values: thin, normal, plump)
 - $a_3 \equiv \text{likesSimon}$ (possible values: yes, no)
 - $a_4 \equiv \text{Pose}$ (possible values: arrogant, natural, goofy)
 - $a_5 \equiv \text{Smile}$ (possible values: none, pleasant, toothy)
 - $a_6 \equiv \text{Smart}$ (possible values: yes, no)
 - target f-on representation: $h \equiv c': \langle a_1, a_2, a_3, a_4, a_5, a_6 \rangle \rightarrow \{0, 1\}$
 - training examples D : positive and negative examples of target function c
- **Aim:** Find a hypothesis $h \in H$ such that $(\forall d \in D) h(d) - c(d) < \epsilon$, where H is the set of all possible hypotheses $h \equiv \langle a_1, a_2, a_3, a_4, a_5, a_6 \rangle$, where each $a_k, k = [1..6]$, may be '?' (\equiv any value is acceptable), '0' (\equiv no value is acceptable), or a specific value.

$h \equiv \langle ?, ?, ?, ?, ?, ? \rangle$ $h \equiv \langle 0, 0, 0, 0, 0, 0 \rangle$ $h \equiv \langle ?, ?, \text{yes}, ?, ?, ? \rangle$

Concept Learning as Search

- Concept learning task:
 - target concept: Girls who Simon likes
 - target function: $c: D \rightarrow \{0, 1\}$
 - data $d \in D$: Girls, each described in terms of the following attributes
 - $a_1 \equiv \text{Hair}$ (possible values: blond, brown, black)
 - $a_2 \equiv \text{Body}$ (possible values: thin, normal, plump)
 - $a_3 \equiv \text{likesSimon}$ (possible values: yes, no)
 - $a_4 \equiv \text{Pose}$ (possible values: arrogant, natural, goofy)
 - $a_5 \equiv \text{Smile}$ (possible values: none, pleasant, toothy)
 - $a_6 \equiv \text{Smart}$ (possible values: yes, no)
 - target f-on representation: $h \equiv c': \langle a_1, a_2, a_3, a_4, a_5, a_6 \rangle \rightarrow \{0, 1\}$
 - training examples D : positive and negative examples of target function c

instances

$$|H| = 1 + 4 \cdot 4 \cdot 3 \cdot 4 \cdot 4 \cdot 3 = 2305$$

$$h \equiv \langle 0, 0, 0, 0, 0, 0 \rangle$$

error rate

- **Aim:** Find a hypothesis $h \in H$ such that $(\forall d \in D) h(d) - c(d) < \epsilon \approx 0$, where H is the set of all possible hypotheses $h \equiv \langle a_1, a_2, a_3, a_4, a_5, a_6 \rangle$, where each a_k , $k = [1..6]$, may be '?' (\equiv any value is acceptable), '0' (\equiv no value is acceptable), or a specific value.

concept learning \equiv searching through H

General-to-Specific Ordering

- Many concept learning algorithms utilize general-to-specific ordering of hypotheses
- General-to-Specific Ordering:
 - $h1$ precedes (is more general than) $h2 \Leftrightarrow (\forall d \in D) (h1(d) = 1) \leftarrow (h2(d) = 1)$
(e.g., $h1 \equiv \langle ?, ?, yes, ?, ?, ? \rangle$ and $h2 \equiv \langle ?, ?, yes, ?, ?, yes \rangle \Rightarrow h1 >_g h2$)
 - $h1$ and $h2$ are of equal generality \Leftrightarrow
 $(\exists d \in D) \{ [(h1(d) = 1) \rightarrow (h2(d) = 1)] \wedge [(h2(d) = 1) \rightarrow (h1(d) = 1)] \}$
(e.g., $h1 \equiv \langle ?, ?, yes, ?, ?, ? \rangle$ and $h2 \equiv \langle ?, ?, ?, ?, ?, yes \rangle \Rightarrow h1 =_g h2$)
 - $h2$ succeeds (is more specific than) $h1 \Leftrightarrow (\forall d \in D) (h1(d) = 1) \leftarrow (h2(d) = 1)$
(e.g., $h1 \equiv \langle ?, ?, yes, ?, ?, ? \rangle$ and $h2 \equiv \langle ?, ?, yes, ?, ?, yes \rangle \Rightarrow h2 \not\geq_g h1$)

Find-S Algorithm

1. Initialise $h \in H$ to the most specific hypothesis: $h \leftarrow \langle a_1, \dots, a_n \rangle, (\forall i) a_i = 0$.
2. FOR each positive training instance $d \in D$, do:
 FOR each attribute $a_i, i = [1..n]$, in h , do:
 IF a_i is satisfied by d
 THEN do nothing
 ELSE replace a_i in h so that the resulting $h' >_g h, h \leftarrow h'$.
3. Output hypothesis h .

Find-S Algorithm – Example

1. Initialise $h \in H$ to the most specific hypothesis: $h \leftarrow \langle a_1, \dots, a_n \rangle, (\forall i) a_i = 0$.
2. FOR each positive training instance $d \in D$, do:
 - FOR each attribute $a_i, i = [1..n]$, in h , do:
 - IF a_i is satisfied by d
 - THEN do nothing
 - ELSE replace a_i in h so that the resulting $h' >_g h, h \leftarrow h'$.
3. Output hypothesis h .

	$c(d)$	<i>hair</i>	<i>body</i>	<i>likesSimon</i>	<i>pose</i>	<i>smile</i>	<i>smart</i>
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

$h \leftarrow \langle 0, 0, 0, 0, 0, 0 \rangle \rightarrow h \equiv d1 \rightarrow h \leftarrow \langle \text{blond}, ?, \text{yes}, ?, ?, \text{no} \rangle$

Find-S Algorithm

- Find-S is guaranteed to output the most specific hypothesis h that best fits positive training examples.
- The hypothesis h returned by Find-S will also fit negative examples as long as training examples are correct.
- However,
 - Find-S is sensitive to noise that is (almost always) present in training examples.
 - there is no guarantee that h returned by Find-S is the *only* h that fits the data.
 - several maximally specific hypotheses may exist that fits the data but, Find-S will output only one.
 - Why we should prefer most specific hypotheses over, e.g., most general hypotheses?

Find-S Algorithm – Example

1. Initialise $h \in H$ to the most specific hypothesis: $h \leftarrow \langle a_1, \dots, a_n \rangle, (\forall i) a_i = 0$.
2. FOR each positive training instance $d \in D$, do:
 - FOR each attribute $a_i, i = [1..n]$, in h , do:
 - IF a_i is satisfied by d
 - THEN do nothing
 - ELSE replace a_i in h so that the resulting $h' >_g h, h \leftarrow h'$.
3. Output hypothesis h .

	$c(d)$	$hair$	$body$	$likesSimon$	$pose$	$smile$	$smart$
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

Find-S $\rightarrow h = \langle \text{blond}, ?, \text{yes}, ?, ?, \text{no} \rangle$ BUT $h_2 = \langle \text{blond}, ?, ?, ?, ?, \text{no} \rangle$ fits D as well

Find-S Algorithm

- Find-S is guaranteed to output the most specific hypothesis h that best fits positive training examples.
- The hypothesis h returned by Find-S will also fit negative examples as long as training examples are correct.
- However,
 - Find-S is sensitive to noise that is (almost always) present in training examples.
 - there is no guarantee that h returned by Find-S is the *only* h that fits the data.
 - several maximally specific hypotheses may exist that fits the data but, Find-S will output only one.
 - Why we should prefer most specific hypotheses over, e.g., most general hypotheses?

Find-S Algorithm – Example

1. Initialise $h \in H$ to the most specific hypothesis: $h \leftarrow \langle a_1, \dots, a_n \rangle, (\forall i) a_i = 0$.
2. FOR each positive training instance $d \in D$, do:
 - FOR each attribute $a_i, i = [1..n]$, in h , do:
 - IF a_i is satisfied by d
 - THEN do nothing
 - ELSE replace a_i in h so that the resulting $h' >_g h, h \leftarrow h'$.
3. Output hypothesis h .

	$c(d)$	$hair$	$body$	$likesSimon$	$pose$	$smile$	$smart$
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

Find-S $\rightarrow h_1 = \langle \text{blond}, ?, ?, ?, ?, \text{no} \rangle$ YET $h_2 = \langle \text{blond}, ?, \text{yes}, ?, ?, ? \rangle$ fits D as well

Find-S Algorithm

- Find-S is guaranteed to output the most specific hypothesis h that best fits positive training examples.
- The hypothesis h returned by Find-S will also fit negative examples as long as training examples are correct.
- However,
 1. Find-S is sensitive to noise that is (almost always) present in training examples.
 2. there is no guarantee that h returned by Find-S is the *only* h that fits the data.
 3. several maximally specific hypotheses may exist that fits the data but, Find-S will output only one.
 4. Why we should prefer most specific hypotheses over, e.g., most general hypotheses?

Candidate-Elimination Algorithm

- Find-S is guaranteed to output the most specific hypothesis h that best fits positive training examples.
- The hypothesis h returned by Find-S will also fit negative examples as long as training examples are correct.
- However,
 1. Find-S is sensitive to noise that is (almost always) present in training examples.
 2. there is no guarantee that h returned by Find-S is the *only* h that fits the data.
 3. several maximally specific hypotheses may exist that fits the data but, Find-S will output only one.
 4. Why we should prefer most specific hypotheses over, e.g., most general hypotheses?

To address the last three drawbacks of Find-S, Candidate-Elimination was proposed

Candidate-Elimination (C-E) Algorithm

- Main idea: Output a set of hypothesis $VS \subseteq H$ that fit (are consistent) with data D
- Candidate-Elimination (C-E) Algorithm is based upon:
 - general-to-specific ordering of hypotheses
 - Def: h is consistent (fits) data $D \Leftrightarrow (\forall \langle d, c(d) \rangle) h(d) = c(d)$
 - Def: version space $VS \subseteq H$ is set of all $h \in H$ that are consistent with D
- C-E algorithm defines VS in terms of two boundaries:
 - general boundary $G \subseteq VS$ is a set of all $h \in VS$ that are the most general
 - specific boundary $S \subseteq VS$ is a set of all $h \in VS$ that are the most specific

Candidate-Elimination (C-E) Algorithm

1. Initialise $G \in VS$ to the most general hypothesis: $h \leftarrow \langle a_1, \dots, a_n \rangle, (\forall i) a_i = ?$.
Initialise $S \in VS$ to the most specific hypothesis: $h \leftarrow \langle a_1, \dots, a_n \rangle, (\forall i) a_i = 0$.
2. FOR each training instance $d \in D$, do:
 - IF d is a positive example
 - Remove from G all h that are not consistent with d .
 - FOR each hypothesis $s \in S$ that is not consistent with d , do:
 - replace s with all h that are consistent with d , $h >_g s$, $\nexists g \in G$,
 - remove from S all s being more general than other s in S .
 - IF d is a negative example
 - Remove from S all h that are not consistent with d .
 - FOR each hypothesis $g \in G$ that is not consistent with d , do:
 - replace g with all h that are consistent with d , $g >_g h$, $h >_g s \in S$,
 - remove from G all g being less general than other g in G .
3. Output hypothesis G and S .

C-E Algorithm – Example

	<i>c(d)</i>	<i>hair</i>	<i>body</i>	<i>likesSimon</i>	<i>pose</i>	<i>smile</i>	<i>smart</i>
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

$G_0 \leftarrow \{\langle ?, ?, ?, ?, ?, ? \rangle\}$, $S_0 \leftarrow \{\langle 0, 0, 0, 0, 0, 0 \rangle\}$

C-E Algorithm – Example

	$c(d)$	<i>hair</i>	<i>body</i>	<i>likesSimon</i>	<i>pose</i>	<i>smile</i>	<i>smart</i>
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

dl is positive \rightarrow *refine S*

no $g \in G_0$ is inconsistent with dl $\rightarrow G_1 \leftarrow G_0 \equiv \{\langle ?, ?, ?, ?, ?, ? \rangle\}$

add to S all minimal generalizations of $s \in S_0$ such that $s \in S_1$ is consistent with dl

$S_1 \leftarrow \{\langle \text{blond, thin, yes, arrogant, toothy, no} \rangle\}$

C-E Algorithm – Example

	$c(d)$	<i>hair</i>	<i>body</i>	<i>likesSimon</i>	<i>pose</i>	<i>smile</i>	<i>smart</i>
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

$d2$ is negative \rightarrow refine G

no $s \in S_1$ is inconsistent with $d2 \rightarrow S_2 \leftarrow S_1 \equiv \{\langle \text{blond, thin, yes, arrogant, toothy, no} \rangle\}$

add to G all minimal specializations of $g \in G_1$ such that $g \in G_2$ is consistent with $d2$

$G_1 \equiv \{\langle ?, ?, ?, ?, ?, ? \rangle\}$

$G_2 \leftarrow \{\langle \text{blond, ?, ?, ?, ?, ?} \rangle, \langle ?, ?, \text{yes, ?, ?, ?} \rangle, \langle ?, ?, ?, \text{arrogant, ?, ?} \rangle, \langle ?, ?, ?, ?, \text{toothy, ?} \rangle, \langle ?, ?, ?, ?, ?, \text{no} \rangle\}$

C-E Algorithm – Example

	$c(d)$	<i>hair</i>	<i>body</i>	<i>likesSimon</i>	<i>pose</i>	<i>smile</i>	<i>smart</i>
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

$d3$ is positive \rightarrow refine S

two $g \in G_2$ are inconsistent with $d3$, i.e., $\langle ?, ?, ?, \text{arrogant}, ?, ? \rangle$ and $\langle ?, ?, ?, ?, \text{toothy}, ? \rangle \rightarrow$
 $G_3 \leftarrow \{ \langle \text{blond}, ?, ?, ?, ?, ? \rangle, \langle ?, ?, \text{yes}, ?, ?, ? \rangle, \langle ?, ?, ?, ?, ?, \text{no} \rangle \}$

add to S all minimal generalizations of $s \in S_2$ such that $s \in S_3$ is consistent with $d3$

$S_2 \equiv \{ \langle \text{blond}, \text{thin}, \text{yes}, \text{arrogant}, \text{toothy}, \text{no} \rangle \}$

$S_3 \leftarrow \{ \langle \text{blond}, ?, \text{yes}, ?, ?, \text{no} \rangle \}$

C-E Algorithm – Example

	$c(d)$	<i>hair</i>	<i>body</i>	<i>likesSimon</i>	<i>pose</i>	<i>smile</i>	<i>smart</i>
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

$d4$ is negative \rightarrow refine G

no $s \in S_3$ is inconsistent with $d4 \rightarrow S_4 \leftarrow S_3 \equiv \{\langle \text{blond}, ?, \text{yes}, ?, ?, \text{no} \rangle\}$

add to G all minimal specializations of $g \in G_3$ such that $g \in G_4$ is consistent with $d4$

$G_3 \equiv \{\langle \text{blond}, ?, ?, ?, ?, ? \rangle, \langle ?, ?, \text{yes}, ?, ?, ? \rangle, \langle ?, ?, ?, ?, ?, \text{no} \rangle\}$

$G_4 \leftarrow \{\langle \text{blond}, ?, ?, ?, ?, ? \rangle, \langle ?, ?, \text{yes}, ?, ?, ? \rangle\}$

C-E Algorithm – Example

	$c(d)$	<i>hair</i>	<i>body</i>	<i>likesSimon</i>	<i>pose</i>	<i>smile</i>	<i>smart</i>
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

$d5$ is negative \rightarrow refine G

no $s \in S_4$ is inconsistent with $d4 \rightarrow S_5 \leftarrow S_4 \equiv \{\langle \text{blond}, ?, \text{yes}, ?, ?, \text{no} \rangle\}$

add to G all minimal specializations of $g \in G_4$ such that $g \in G_5$ is consistent with $d5$

$G_4 \equiv \{\langle \text{blond}, ?, ?, ?, ?, ? \rangle, \langle ?, ?, \text{yes}, ?, ?, ? \rangle\}$

$G_5 \leftarrow \{\langle \text{blond}, \text{X}?, ?, ?, \text{no} \rangle, \langle ?, ?, \text{yes}, ?, ?, ? \rangle\}$

C-E Algorithm – Example

	$c(d)$	<i>hair</i>	<i>body</i>	<i>likesSimon</i>	<i>pose</i>	<i>smile</i>	<i>smart</i>
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

Output of C-E:

*version space of hypotheses $VS \subseteq H$ bound with
specific boundary $S \equiv \{\langle \text{blond}, ?, \text{yes}, ?, ?, \text{no} \rangle\}$ and
general boundary $G \equiv \{\langle ?, ?, \text{yes}, ?, ?, ? \rangle\}$*

Output of Find-S:

most specific hypothesis $h \equiv \langle \text{blond}, ?, \text{yes}, ?, ?, \text{no} \rangle$

C-E Algorithm – Example

	$c(d)$	<i>hair</i>	<i>body</i>	<i>likesSimon</i>	<i>pose</i>	<i>smile</i>	<i>smart</i>
1	1	blond	thin	yes	arrogant	toothy	no
2	0	brown	thin	no	natural	pleasant	yes
3	1	blond	plump	yes	goofy	pleasant	no
4	0	black	thin	no	arrogant	none	no
5	0	blond	plump	no	natural	toothy	yes

Output of C-E:

*version space of hypotheses $VS \subseteq H$ bound with
specific boundary $S \equiv \{\langle \text{blond}, ?, \text{yes}, ?, ?, \text{no} \rangle\}$ and
general boundary $G \equiv \{\langle ?, ?, \text{yes}, ?, ?, ? \rangle\}$*

$VS \equiv \{\langle ?, ?, \text{yes}, ?, ?, ? \rangle, \langle \text{blond}, ?, \text{yes}, ?, ?, ? \rangle, \langle ?, ?, \text{yes}, ?, ?, \text{no} \rangle, \langle \text{blond}, ?, \text{yes}, ?, ?, \text{no} \rangle\}$

Concept Learning – Lecture Overview

- Why machine learning?
- Well-posed learning problems
- Designing a machine learning system
- Concept learning task
- Concept learning as Search
- Find-S algorithm
- Candidate-Elimination algorithm

Concept Learning – Exam Questions

- Tom Mitchell's book – chapter 1 and chapter 2
- Relevant exercises from chapter 1: 1.1, 1.2, 1.3, 1.5
- Relevant exercises from chapter 2: 2.1, 2.2, 2.3, 2.4, 2.5

Course 395: Machine Learning – Lectures

- Lecture 1-2: Concept Learning (*M. Pantic*)
- Lecture 3-4: Decision Trees & CBC Intro (*M. Pantic*)
- Lecture 5-6: Artificial Neural Networks (*THs*)
- Lecture 7-8: Instance Based Learning (*M. Pantic*)
- Lecture 9-10: Genetic Algorithms (*M. Pantic*)
- Lecture 11-12: Evaluating Hypotheses (*THs*)
- Lecture 13-14: Guest Lectures on ML Applications
- Lecture 15-16: Inductive Logic Programming (*S. Muggleton*)
- Lecture 17-18: Inductive Logic Programming (*S. Muggleton*)